

# heart

December 1, 2023

Importing the Dependencies

```
[ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

Data Collection and Processing

```
[ ]: # loading the csv data to a Pandas DataFrame
heart_data = pd.read_csv('/content/heart.csv')
```

```
[ ]: # print first 5 rows of the dataset
heart_data.head()
```

```
[ ]:
   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  \
0   63   1   3     145    233   1         0     150      0        2.3      0
1   37   1   2     130    250   0         1     187      0        3.5      0
2   41   0   1     130    204   0         0     172      0        1.4      2
3   56   1   1     120    236   0         1     178      0        0.8      2
4   57   0   0     120    354   0         1     163      1        0.6      2

   ca  thal  target
0   0     1        1
1   0     2        1
2   0     2        1
3   0     2        1
4   0     2        1
```

```
[ ]: # print last 5 rows of the dataset
heart_data.tail()
```

```
[ ]:
   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
298   57   0   0     140    241   0         1     123      1        0.2
299   45   1   3     110    264   0         1     132      0        1.2
300   68   1   0     144    193   1         1     141      0        3.4
301   57   1   0     130    131   0         1     115      1        1.2
```

302	57	0	1	130	236	0	0	174	0	0.0
-----	----	---	---	-----	-----	---	---	-----	---	-----

	slope	ca	thal	target
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

```
[ ]: # number of rows and columns in the dataset
heart_data.shape
```

```
[ ]: (303, 14)
```

```
[ ]: # getting some info about the data
heart_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null    int64
 1   sex         303 non-null    int64
 2   cp          303 non-null    int64
 3   trestbps    303 non-null    int64
 4   chol        303 non-null    int64
 5   fbs         303 non-null    int64
 6   restecg     303 non-null    int64
 7   thalach     303 non-null    int64
 8   exang       303 non-null    int64
 9   oldpeak     303 non-null    float64
10   slope       303 non-null    int64
11   ca          303 non-null    int64
12   thal        303 non-null    int64
13   target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
[ ]: # checking for missing values
heart_data.isnull().sum()
```

```
[ ]: age         0
sex         0
cp          0
trestbps    0
chol        0
```

```

fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64

```

```
[ ]: # statistical measures about the data
heart_data.describe()
```

```
[ ]:
count    age      sex      cp      trestbps      chol      fbs  \
mean     54.366337  0.683168  0.966997  131.623762  246.264026  0.148515
std       9.082101  0.466011  1.032052  17.538143   51.830751  0.356198
min      29.000000  0.000000  0.000000   94.000000  126.000000  0.000000
25%      47.500000  0.000000  0.000000  120.000000  211.000000  0.000000
50%      55.000000  1.000000  1.000000  130.000000  240.000000  0.000000
75%      61.000000  1.000000  2.000000  140.000000  274.500000  0.000000
max      77.000000  1.000000  3.000000  200.000000  564.000000  1.000000

count    restecg    thalach    exang    oldpeak    slope    ca  \
mean      0.528053  149.646865  0.326733  1.039604   1.399340  0.729373
std       0.525860  22.905161  0.469794  1.161075   0.616226  1.022606
min       0.000000  71.000000  0.000000  0.000000   0.000000  0.000000
25%       0.000000  133.500000  0.000000  0.000000   1.000000  0.000000
50%       1.000000  153.000000  0.000000  0.800000   1.000000  0.000000
75%       1.000000  166.000000  1.000000  1.600000   2.000000  1.000000
max       2.000000  202.000000  1.000000  6.200000   2.000000  4.000000

count    thal    target
mean     2.313531  0.544554
std      0.612277  0.498835
min      0.000000  0.000000
25%      2.000000  0.000000
50%      2.000000  1.000000
75%      3.000000  1.000000
max      3.000000  1.000000

```

```
[ ]: # checking the distribution of Target Variable
heart_data['target'].value_counts()
```

```
[ ]: 1    165
      0    138
      Name: target, dtype: int64
```

1 -> Defective Heart

0 -> Healthy Heart

```
[ ]: # To find the correlation in data set we use corr()
import matplotlib.pyplot as plt
import seaborn as sns
corrMatrix=heart_data.corr()
corrMatrix
```


```
[ ]:
      age      sex      cp  trestbps      chol      fbs  \
age      1.000000 -0.098447 -0.068653  0.279351  0.213678  0.121308
sex     -0.098447  1.000000 -0.049353 -0.056769 -0.197912  0.045032
cp      -0.068653 -0.049353  1.000000  0.047608 -0.076904  0.094444
trestbps 0.279351 -0.056769  0.047608  1.000000  0.123174  0.177531
chol     0.213678 -0.197912 -0.076904  0.123174  1.000000  0.013294
fbs      0.121308  0.045032  0.094444  0.177531  0.013294  1.000000
restecg  -0.116211 -0.058196  0.044421 -0.114103 -0.151040 -0.084189
thalach  -0.398522 -0.044020  0.295762 -0.046698 -0.009940 -0.008567
exang     0.096801  0.141664 -0.394280  0.067616  0.067023  0.025665
oldpeak   0.210013  0.096093 -0.149230  0.193216  0.053952  0.005747
slope    -0.168814 -0.030711  0.119717 -0.121475 -0.004038 -0.059894
ca        0.276326  0.118261 -0.181053  0.101389  0.070511  0.137979
thal      0.068001  0.210041 -0.161736  0.062210  0.098803 -0.032019
target   -0.225439 -0.280937  0.433798 -0.144931 -0.085239 -0.028046

      restecg  thalach  exang  oldpeak  slope  ca  \
age     -0.116211 -0.398522  0.096801  0.210013 -0.168814  0.276326
sex     -0.058196 -0.044020  0.141664  0.096093 -0.030711  0.118261
cp       0.044421  0.295762 -0.394280 -0.149230  0.119717 -0.181053
trestbps -0.114103 -0.046698  0.067616  0.193216 -0.121475  0.101389
chol     -0.151040 -0.009940  0.067023  0.053952 -0.004038  0.070511
fbs      -0.084189 -0.008567  0.025665  0.005747 -0.059894  0.137979
restecg   1.000000  0.044123 -0.070733 -0.058770  0.093045 -0.072042
thalach   0.044123  1.000000 -0.378812 -0.344187  0.386784 -0.213177
exang    -0.070733 -0.378812  1.000000  0.288223 -0.257748  0.115739
oldpeak  -0.058770 -0.344187  0.288223  1.000000 -0.577537  0.222682
slope     0.093045  0.386784 -0.257748 -0.577537  1.000000 -0.080155
ca       -0.072042 -0.213177  0.115739  0.222682 -0.080155  1.000000
thal     -0.011981 -0.096439  0.206754  0.210244 -0.104764  0.151832
target    0.137230  0.421741 -0.436757 -0.430696  0.345877 -0.391724

      thal  target
age      0.068001 -0.225439
```

sex	0.210041	-0.280937
cp	-0.161736	0.433798
trestbps	0.062210	-0.144931
chol	0.098803	-0.085239
fbs	-0.032019	-0.028046
restecg	-0.011981	0.137230
thalach	-0.096439	0.421741
exang	0.206754	-0.436757
oldpeak	0.210244	-0.430696
slope	-0.104764	0.345877
ca	0.151832	-0.391724
thal	1.000000	-0.344029
target	-0.344029	1.000000

```
[ ]: import seaborn as sns
sns.heatmap(corrMatrix,cmap="YlGnBu",annot=True)
plt.gcf().set_size_inches(10, 10)
```



Splitting the Features and Target

```
[ ]: X = heart_data.drop(columns='target', axis=1)
     Y = heart_data['target']
```

```
[ ]: print(X)
```

```

    age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0    63   1   3     145    233   1         0     150     0        2.3
1    37   1   2     130    250   0         1     187     0        3.5
2    41   0   1     130    204   0         0     172     0        1.4
3    56   1   1     120    236   0         1     178     0        0.8
```

4	57	0	0	120	354	0	1	163	1	0.6
..	...	..	...	...	...	...	...	...	...	...
298	57	0	0	140	241	0	1	123	1	0.2
299	45	1	3	110	264	0	1	132	0	1.2
300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal
0	0	0	1
1	0	0	2
2	2	0	2
3	2	0	2
4	2	0	2
..	...	..	...
298	1	0	3
299	1	0	3
300	1	2	3
301	1	1	3
302	1	1	2

[303 rows x 13 columns]

```
[ ]: print(Y)
```

```
0      1
1      1
2      1
3      1
4      1
..
298    0
299    0
300    0
301    0
302    0
```

Name: target, Length: 303, dtype: int64

Splitting the Data into Training data & Test Data

```
[ ]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
↳stratify=Y, random_state=2)
```

```
[ ]: print(X.shape, X_train.shape, X_test.shape)
```

```
(303, 13) (242, 13) (61, 13)
```

Model Training

Random Forest

```
[ ]: from sklearn.ensemble import RandomForestClassifier
model=RandomForestClassifier(n_estimators=10)
```

```
[ ]: # training the LogisticRegression model with Training data
model.fit(X_train, Y_train)
```

```
[ ]: RandomForestClassifier(n_estimators=10)
```

Model Evaluation

Accuracy Score

```
[ ]: # accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
[ ]: print('Accuracy on Training data : ', training_data_accuracy)
```

Accuracy on Training data : 0.987603305785124

```
[ ]: # accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
[ ]: print('Accuracy on Test data : ', test_data_accuracy)
```

Accuracy on Test data : 0.7540983606557377

Building a Predictive System

```
[ ]: input_data = (62,0,0,140,268,0,0,160,0,3.6,0,2,2)

# change the input data to a numpy array
input_data_as_numpy_array= np.asarray(input_data)

# reshape the numpy array as we are predicting for only on instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')
```

[0]

The Person does not have a Heart Disease



```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names
```

```
warnings.warn(
```

Saving the trained model

```
[ ]: import pickle
```

```
[ ]: filename = 'heart_disease_model.sav'
pickle.dump(model, open(filename, 'wb'))
```

```
[ ]: # loading the saved model
loaded_model = pickle.load(open('heart_disease_model.sav', 'rb'))
```

```
[ ]: for column in X.columns:
    print(column)
```

```
age
sex
cp
trestbps
chol
fbs
restecg
thalach
exang
oldpeak
slope
ca
thal
```

```
[ ]:
```